

 Search models, datasets, users...[Terms of Service](#)[Privacy Policy](#)[Content Policy](#)[Code of Conduct](#)

# Content Policy



Effective Date: April 10, 2025

At Hugging Face, we are dedicated to helping the entire AI Community work together to advance **Open, Collaborative, and Responsible** Machine Learning. We value these relationships and aim to **maintain an environment where people feel welcome** and supported and get the most out of their contributions and experiences.

We created the following Content Policy to outline our actions to **protect our Community** on our Platform and to provide important information about the **Content we authorize** to be posted on our Platform.

This Policy is incorporated into our Terms of Service, which is a **binding agreement** between you and us. You should carefully review all of our other guidelines, policies, and documents available on our Website, including our [Terms of Service](#) and [Privacy Policy](#). Also, read the Hugging Face [Code of Conduct](#), which gives a general overview and states our standards and how we wish the community would behave.

**By accessing** or using our Platform, **you consent** to all of this Policy and our other policies and terms.

So, if you do not agree with any of those, please do not access or use our Platform. We **may change** or update this Policy from time to time. Changes will be effective 10 days following posting the updated Policy on the Website. If you continue using our Platform 10 days following such posting, that means you accept those changes.



# Key definitions

Capitalized terms used but not defined herein shall have the meaning assigned in our [Terms of Service](#), [Privacy Policy](#), and all other policies available on our Website.

“**Content**” refers to any material posted, displayed, or accessed on our Website or Hub, including but not limited to code, data, text, graphics, images, username, applications, or software you, we, or any third party provide or make available.

Content types may include:

- “**ML Artifacts**”: Code and assets hosted as Hugging Face Repositories, including Models, Datasets, Spaces;
- “**Community Content**”: Content that can be found in the Community section of the Hugging Face Platform, including discussions, comments, and usernames, as well as related documentation such as READMEs, model cards, data cards, pull requests, and merges.

“**Community**” refers to all Users of the Hugging Face Platform, including Hugging Face personnel.

“**Community Tab**” refers to a collaborative feature where the Community can discuss specific Repositories, including providing feedback, brainstorming ideas and opening pull requests for improvements.

“**Hugging Face**” refers to Hugging Face Inc., which may perform its obligations through its affiliates, directors, subsidiaries, contractors, licensors, officers, agents, and/or employees.

“**Platform**”, “**Hugging Face Hub**”, or “**Hub**” refers to the hosting platform available at [huggingface.co](https://huggingface.co) where Users can build, benchmark, share, version, and deploy Repositories, which may include Models, Datasets, and Machine Learning Applications.

“**Repository**” refers to a data structure that contains all of the project files and the entire revision history.

A Repository may be:

- “**Public**”: anyone on the internet can see it, but only you or members of your organization can make changes;
- “**Private**”: only you or members of your organization can see and make changes to the Repository; New Users need to join the maintaining organization in order to both see the Repository and access its Content.
- “**Gated**”: Gated Repositories and their Community Content are visible to everyone, but access to their ML artifacts (data, model weights) requires either accepting conditions in a click-through form or approval by the Repository maintainers.
- “**Disabled**”: a Repository that has its access blocked to all Community members except its owner.

“**Repository tag**” refers to a way to associate and present metadata of any kind related to a repository.

“**Team**” or “**Hugging Face Team**” refers to Hugging Face personnel.

## **Restricted Content**

**Restricted Content** includes clear violations of our **Content Policy** or **applicable laws**, and is subject to immediate action.

If you see Content that violates this Policy, **report it**. It will be directly addressed by the Hugging Face Team on a case-by-case basis.

Restricted Content may include:

### **1. Unlawful or illegal Content**

- All Content that violates any applicable law or regulation.
- Content promoting high-risk illegal activities (weapons development, illegal substances, scams, gambling, pseudo-pharmaceuticals, plagiarism, etc.).

## 2. Fraudulent or Malicious Activities

- Unlawful, defamatory, fraudulent, or intentionally deceptive Content (e.g., disinformation, phishing, scams, inauthentic behavior).
- Content that promotes or induces unlawful or fraudulent currencies, securities, investments, or transactions.
- Content that performs illegal or unlicensed professional practices (medical, legal, financial, etc.).

## 3. Harmful or Abusive Content

- Content that harms individuals or groups.
- Content promoting discrimination or hate speech (see our Code of Conduct).
- Content involving harassment, bullying, or demeaning behavior.
- Sexual Content used for harassment, bullying, or created without explicit consent.
- Underage nudity or any sexual Content involving minors.
- Terrorist Content or Content that glorifies violence, suffering, or humiliation.

## 4. Privacy and Intellectual Property (IP)

- Content that violates the privacy of a third party, including but not limited to publishing others' private information, such as a physical or email address, without their explicit permission.
- Content that infringes the intellectual property rights of a third party.

## 5. Platform Abuse, Security Violations and Spam

- Content designed to disrupt, damage, or gain unauthorized access to systems or devices.
- Content that attempts to transmit or generate malicious code (e.g., malware, trojans, viruses).
- Abuse or interference with Hugging Face services, including:
  - Using unauthorized bot APIs or remote management tools.
  - Hosting excessive or irrelevant data in repositories.
  - Using tools like Cloudflare Tunnel, TOR, proxies, VNC, Chrome Remote Server, etc., to bypass restrictions.
  - Incentivizing manipulation of Hugging Face Hub metrics (e.g., exchanging rewards for likes).
  - Cryptomining practices.
  - Spam (e.g., advertising products/services, excessive bulk activity, or disrupting user experience).

While the categories of Restricted Content above provide a clear framework, we may also moderate other types of Content in response to evolving challenges posed by advancements in Machine Learning. As we assess such Content, we hold consent as a core value, ensuring our approach remains thoughtful, adaptive, and respectful of individual and community rights.

Users see the same public content on the Hub without personalized recommendations. Trending content is influenced by the number of likes in the past few days. Posts and updates appear from accounts users actively follow, displayed in strict chronological order without ranking or algorithmic curation.



## How you can report Content

## 1. Any Content

If you encounter Content on our Platform that you believe violates this Policy, you can report it to us through the following channels:

- In-Platform Reporting: Click the  button and select the "Report" option. **This will:**
  - **For Repositories:** Open a public report and notify the Hugging Face Team
  - **For Posts and Comments:** Notify the Hugging Face Team directly
- Email: Send a report to [safety@huggingface.co](mailto:safety@huggingface.co)

Please note that **reports are Community Content** and are themselves subject to this Policy. Abusive uses of the flagging feature, including but not limited to spamming or harassment, will not be tolerated.

In addition, in some situations, the Hugging Face Team may flag Content to monitor requests or concerns expressed through other channels (e.g., via social media, automatic detection, or raised internally). Depending on the severity of the issue, we may take action as described in the Section below, “How we moderate Content”.

## 2. IP Infringement and DMCA takedown process

If you believe that any Content on our website **infringes** upon your **intellectual property rights**, you can submit a Takedown notice to [dmca@huggingface.co](mailto:dmca@huggingface.co). This is a process we follow according to applicable law.

Please include detailed and accurate information to support your claim. By submitting a complaint, you confirm that all information provided is truthful and that you will not intentionally submit any false or misleading details.

We will review the takedown notice to ensure it meets legal requirements and isn't abusive. If this is confirmed, we will inform the User who uploaded the allegedly infringing Content and disable it. The original uploader can challenge this decision if they think there has been a mistake by submitting a counter-notification to us and similarly including detailed and

accurate information to support their claims. We will review the counter-notification to ensure it meets the legal requirements, and, if it does, we will inform the original claimant, who will have 14 U.S. business days to take legal action to prevent the Content from being restored.

## 🔗 How we moderate Content

If Content is reported or flagged as potentially problematic, the Hugging Face Team will review it to determine whether it meets the terms and conditions of this Content Policy. This requires an iterative approach to determine whether and under what conditions it may be hosted on the Platform. Repository Reports will, when possible, be addressed in collaboration with both the Repository owner and any concerned party, as part of our decision-making process.

After receiving and reviewing a report, our Team will take action on the Content where appropriate. These actions may include, but are not limited to:

- Asking the relevant User for collaboration or modifications to the Content
- Unranking the Content
- Adding a Not for All Audiences (NFAA) Tag
- Removing or Disabling access to the Content
- Restricting Interactions
- Account Suspension or Termination



We prioritize collaborative solutions for both ML Artifacts and Community Content that involve the owner of the Repository whenever possible, especially in cases where modifications or additional guardrails can help the Content meet this Policy. We offer a range of useful tools on our Platform to support Community Content moderation. This regularly updated list allows the Community to play an active role in moderation efforts.

## Repository owners can:

- Edit discussions and PRs titles
- Hide or edit a community comment
- Tag their Content as being “Not For All Audiences”
- Gate the access to their Repository, allowing the owners to manually review and approve/reject access to their ML artifacts (see the documentation for Models and Datasets).
- Close, lock, or delete discussions and PRs

## Authors of a discussion or a pull request can:

- Edit the discussion or PR title
- Edit their own comments
- Hide their own comments
- Close the discussion or PR

## 🔗 How you can contest a moderation decision

If you disagree with a moderation decision to disable your Content or suspend or terminate your account, you can **send us your complaint** at [legal@huggingface.co](mailto:legal@huggingface.co). Please include a  **arguments** and evidence you may have to support your claims so we can reconsider our decision. We will diligently review your complaint and will respond to you.

If you remain unsatisfied, and in accordance with the **Digital Services Act**, you have the right to approach a certified out-of-court dispute settlement body. We commit to participating in this process in good faith, with the aim of reaching a resolution.

## 🔗 Contact us

**For our Users:** For all communications regarding our Content Policy and how we deal with moderating Content at Hugging Face, please contact us at [feedback@huggingface.co](mailto:feedback@huggingface.co).

- ▶ For EU residents
- ▶ Additional content policy in compliance with Australian Section 13(3)(a) of the DIS Standard (Australian Online Safety)

 System theme

## Company

TOS

Privacy

About

Jobs

## Website

Models

Datasets

Spaces

Pricing

Docs



S

X

x